

Run-time Environments

Lecture 11

Profs. Aiken CS 143 Lecture 11

1

Status

- We have covered the front-end phases
 - Lexical analysis
 - Parsing
 - Semantic analysis
- Next are the back-end phases
 - Optimization
 - Code generation
- We'll do code generation first . . .

Profs. Aiken CS 143 Lecture 11

2

Run-time environments

- Before discussing code generation, we need to understand what we are trying to generate
- There are a number of standard techniques for structuring executable code that are widely used

Profs. Aiken CS 143 Lecture 11

3

Outline

- Management of run-time resources
- Correspondence between
 - static (compile-time) and
 - dynamic (run-time) structures
- Storage organization

Profs. Aiken CS 143 Lecture 11

4


Run-time Resources

- Execution of a program is initially under the control of the operating system
- When a program is invoked:
 - The OS allocates space for the program
 - The code is loaded into part of the space
 - The OS jumps to the entry point (i.e., "main")

Profs. Aiken CS 143 Lecture 11

5

Memory Layout



Profs. Aiken CS 143 Lecture 11

6

Notes

- By tradition, pictures of machine organization have:
 - Low address at the top
 - High address at the bottom
 - Lines delimiting areas for different kinds of data
- These pictures are simplifications
 - E.g., not all memory need be contiguous

Profs. Aiken CS 143 Lecture 11

7

What is Other Space?

- Holds all data for the program
- Other Space = Data Space
- Compiler is responsible for:
 - Generating code
 - Orchestrating use of the data area

Profs. Aiken CS 143 Lecture 11

8

Code Generation Goals

- Two goals:
 - Correctness
 - Speed
- Most complications in code generation come from trying to be fast as well as correct

Profs. Aiken CS 143 Lecture 11

9

Assumptions about Execution

1. Execution is sequential; control moves from one point in a program to another in a well-defined order
2. When a procedure is called, control eventually returns to the point immediately after the call

Do these assumptions always hold?

Profs. Aiken CS 143 Lecture 11

10

Activations

- An invocation of procedure *P* is an *activation* of *P*
- The *lifetime* of an activation of *P* is
 - All the steps to execute *P*
 - Including all the steps in procedures *P* calls

Profs. Aiken CS 143 Lecture 11

11

Lifetimes of Variables

- The *lifetime* of a variable *x* is the portion of execution in which *x* is defined
- Note that
 - Lifetime is a dynamic (run-time) concept
 - Scope is a static concept

Profs. Aiken CS 143 Lecture 11

12

Activation Trees

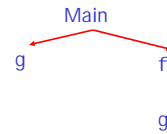
- Assumption (2) requires that when **P** calls **Q**, then **Q** returns before **P** does
- Lifetimes of procedure activations are properly nested
- Activation lifetimes can be depicted as a tree

Profs. Aiken CS 143 Lecture 11

13

Example

```
Class Main {
  g() : Int { 1 };
  f() : Int { g() };
  main() : Int {{ g(); f(); }};
}
```



Profs. Aiken CS 143 Lecture 11

14

Example 2

```
Class Main {
  g() : Int { 1 };
  f(x: Int): Int { if x = 0 then g() else f(x - 1) fi };
  main() : Int {{ f(3); }};
}
```

What is the activation tree for this example?

Profs. Aiken CS 143 Lecture 11

15

Notes

- The activation tree depends on run-time behavior
- The activation tree may be different for every program input
- Since activations are properly nested, a stack can track currently active procedures

Profs. Aiken CS 143 Lecture 11

16

Example

```
Class Main {
  g() : Int { 1 };
  f() : Int { g() };
  main() : Int {{ g(); f(); }};
}
```

Main

Stack
Main

Profs. Aiken CS 143 Lecture 11

17

Example

```
Class Main {
  g() : Int { 1 };
  f() : Int { g() };
  main() : Int {{ g(); f(); }};
}
```



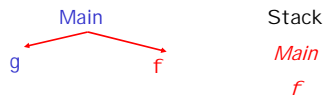
Stack
Main
g

Profs. Aiken CS 143 Lecture 11

18

Example

```
Class Main {
  g() : Int { 1 };
  f(): Int { g() };
  main(): Int {{ g(); f(); }};
}
```

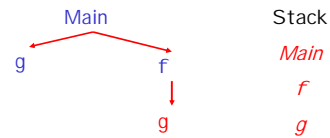


Prof. Aiken CS 143 Lecture 11

19

Example

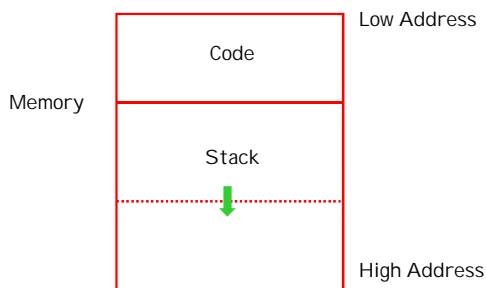
```
Class Main {
  g() : Int { 1 };
  f(): Int { g() };
  main(): Int {{ g(); f(); }};
}
```



Prof. Aiken CS 143 Lecture 11

20

Revised Memory Layout



Prof. Aiken CS 143 Lecture 11

21

Activation Records

- The information needed to manage one procedure activation is called an *activation record (AR)* or *frame*
- If procedure F calls G, then G's activation record contains a mix of info about F and G.

Prof. Aiken CS 143 Lecture 11

22

What is in G's AR when F calls G?

- F is "suspended" until G completes, at which point F resumes. G's AR contains information needed to resume execution of F.
- G's AR may also contain:
 - G's return value (needed by F)
 - Actual parameters to G (supplied by F)
 - Space for G's local variables

Prof. Aiken CS 143 Lecture 11

23

The Contents of a Typical AR for G

- Space for G's return value
- Actual parameters
- Pointer to the previous activation record
 - The *control link*; points to AR of caller of G
- Machine status prior to calling G
 - Contents of registers & program counter
 - Local variables
- Other temporary values

Prof. Aiken CS 143 Lecture 11

24

Example 2, Revisited

```

Class Main {
  g() : Int { 1 };
  f(x: Int): Int { if x=0 then g() else f(x - 1)(**)fi);
  main(): Int {{f(3); (*)
}};}

```

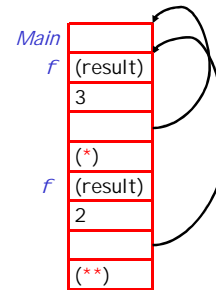
AR for f:

result
argument
control link
return address

Profs. Aiken CS 143 Lecture 11

25

Stack After Two Calls to f



Profs. Aiken CS 143 Lecture 11

26

Notes

- **Main** has no argument or local variables and its result is never used; its AR is uninteresting
- **(*)** and **(**)** are return addresses of the invocations of **f**
 - The return address is where execution resumes after a procedure call finishes
- This is only one of many possible AR designs
 - Would also work for C, Pascal, FORTRAN, etc.

Profs. Aiken CS 143 Lecture 11

27

The Main Point

The compiler must determine, at compile-time, the layout of activation records and generate code that correctly accesses locations in the activation record

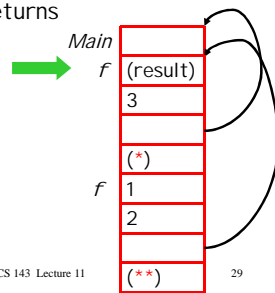
Thus, the AR layout and the code generator must be designed together!

Profs. Aiken CS 143 Lecture 11

28

Example

The picture shows the state after the call to the 2nd invocation of **f** returns



Profs. Aiken CS 143 Lecture 11

29

Discussion

- The advantage of placing the return value 1st in a frame is that the caller can find it at a fixed offset from its own frame
- There is nothing magic about this organization
 - Can rearrange order of frame elements
 - Can divide caller/callee responsibilities differently
 - An organization is better if it improves execution speed or simplifies code generation

Profs. Aiken CS 143 Lecture 11

30

Discussion (Cont.)

- Real compilers hold as much of the frame as possible in registers
 - Especially the method result and arguments

Profs. Aiken CS 143 Lecture 11

31

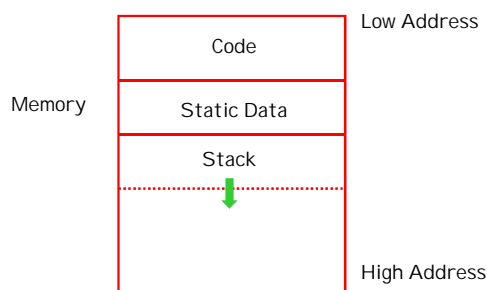
Globals

- All references to a global variable point to the same object
 - Can't store a global in an activation record
- Globals are assigned a fixed address once
 - Variables with fixed address are "statically allocated"
- Depending on the language, there may be other statically allocated values

Profs. Aiken CS 143 Lecture 11

32

Memory Layout with Static Data



Profs. Aiken CS 143 Lecture 11

33

Heap Storage

- A value that outlives the procedure that creates it cannot be kept in the AR

```
method foo() { new Bar }
```

The `Bar` value must survive deallocation of `foo`'s AR
- Languages with dynamically allocated data use a *heap* to store dynamic data

Profs. Aiken CS 143 Lecture 11

34

Notes

- The code area contains object code
 - For most languages, fixed size and read only
- The static area contains data (not code) with fixed addresses (e.g., global data)
 - Fixed size, may be readable or writable
- The stack contains an AR for each currently active procedure
 - Each AR usually fixed size, contains locals
- Heap contains all other data
 - In C, heap is managed by *malloc* and *free*

Profs. Aiken CS 143 Lecture 11

35

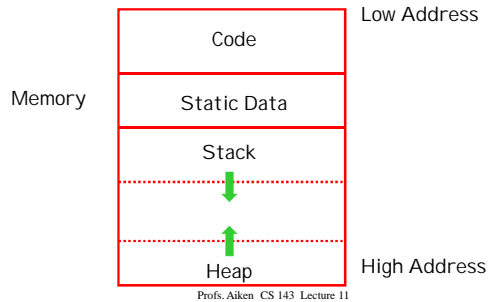
Notes (Cont.)

- Both the heap and the stack grow
- Must take care that they don't grow into each other
- Solution: start heap and stack at opposite ends of memory and let them grow towards each other

Profs. Aiken CS 143 Lecture 11

36

Memory Layout with Heap



Data Layout

- Low-level details of machine architecture are important in laying out data for correct code and maximum performance
- Chief among these concerns is *alignment*

Prof. Aiken CS 143 Lecture 11

38

Alignment

- Most modern machines are (still) 32 bit
 - 8 bits in a byte
 - 4 bytes in a word
 - Machines are either byte or word addressable
- Data is *word aligned* if it begins at a word boundary
- Most machines have some alignment restrictions
 - Or performance penalties for poor alignment

Prof. Aiken CS 143 Lecture 11

39

Alignment (Cont.)

- Example: A string
“Hello”
Takes 5 characters (without a terminating \0)
- To word align next datum, add 3 “padding” characters to the string
- The padding is not part of the string, it’s just unused memory

Prof. Aiken CS 143 Lecture 11

40

Next Topic: Stack Machines

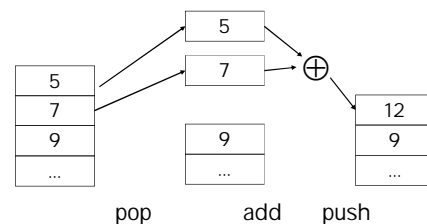
- A simple evaluation model
- No variables or registers
- A stack of values for intermediate results
- Each instruction:
 - Takes its operands from the top of the stack
 - Removes those operands from the stack
 - Computes the required operation on them
 - Pushes the result on the stack

Prof. Aiken CS 143 Lecture 11

41

Example of Stack Machine Operation

- The addition operation on a stack machine



Prof. Aiken CS 143 Lecture 11

42

Example of a Stack Machine Program

- Consider two instructions
 - `push i` - place the integer `i` on top of the stack
 - `add` - pop two elements, add them and put the result back on the stack
- A program to compute $7 + 5$:


```
push 7
push 5
add
```

Prof. Aiken CS 143 Lecture 11

43

Why Use a Stack Machine ?

- Each operation takes operands from the same place and puts results in the same place
- This means a uniform compilation scheme
- And therefore a simpler compiler

Prof. Aiken CS 143 Lecture 11

44

Why Use a Stack Machine ?

- Location of the operands is implicit
 - Always on the top of the stack
- No need to specify operands explicitly
- No need to specify the location of the result
- Instruction `"add"` as opposed to `"add r1, r2"`
 - ⇒ Smaller encoding of instructions
 - ⇒ More compact programs
- This is one reason why Java Bytecodes use a stack evaluation model

Prof. Aiken CS 143 Lecture 11

45

Optimizing the Stack Machine

- The `add` instruction does 3 memory operations
 - Two reads and one write to the stack
 - The top of the stack is frequently accessed
- Idea: keep the top of the stack in a register (called accumulator)
 - Register accesses are faster
- The `"add"` instruction is now


```
acc ← acc + top_of_stack
```

 - Only one memory operation!

Prof. Aiken CS 143 Lecture 11

46

Stack Machine with Accumulator

Invariants

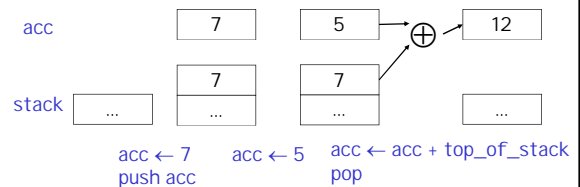
- The result of an expression is in the accumulator
- For `op(e1, ..., en)` push the accumulator on the stack after computing `e1, ..., en-1`
 - After the operation pops `n-1` values
- Expression evaluation preserves the stack

Prof. Aiken CS 143 Lecture 11

47

Stack Machine with Accumulator. Example

- Compute $7 + 5$ using an accumulator



Prof. Aiken CS 143 Lecture 11

48

A Bigger Example: $3 + (7 + 5)$

Code	Acc	Stack
$acc \leftarrow 3$	3	<init>
push acc	3	3, <init>
$acc \leftarrow 7$	7	3, <init>
push acc	7	7, 3, <init>
$acc \leftarrow 5$	5	7, 3, <init>
$acc \leftarrow acc + \text{top_of_stack}$	12	7, 3, <init>
pop	12	3, <init>
$acc \leftarrow acc + \text{top_of_stack}$	15	3, <init>
pop	15	<init>

Profs. Aiken CS 143 Lecture 11

49

Notes

- It is very important evaluation of a subexpression preserves the stack
 - Stack before the evaluation of $7 + 5$ is 3, <init>
 - Stack after the evaluation of $7 + 5$ is 3, <init>
 - The first operand is on top of the stack

Profs. Aiken CS 143 Lecture 11

50